

用 Zebra 在 Linux 上构建路由器实战手册

一、Zebra 介绍

Zebra 是一个 TCP/IP 路由软件，支持 BGP-4、BGP-4+、OSPFv2、OSPFv3、RIPv1、RIPv2 和 RIPv6。它的发行遵循 GNU 通用公共许可协议，可以运行于 Linux 以及其他一些 Unix 变体操作系统上。

Zebra 是那些系统最新的发行版本中的路由软件。最新版本的 Zebra 以及文档可以从 GNU Zebra 网站上下载：<http://www.zebra.org/>

Zebra 的设计独特，采用模块的方法来管理协议。可以根据网络需要启用或者禁用协议。

Zebra 最为实用的一点是它的配置形式同 Cisco IOS 极其类似。尽管它的配置与 IOS 相比还是有一些不同，但是这对于那些已经熟悉 IOS 的网络工程师来说在这种环境下工作将相当自如。

二、安装

1、下载

<ftp://ftp.zebra.org/pub/zebra/zebra-0.95a.tar.gz>

2、安装过程

```
# gunzip zebra-0.95a.tar.gz
```

```
#tar xvf zebra-0.95a.tar
```

会在当前目录下自动生成一个 zebra-0.95a 子目录，进入

zebra-0.95a 目录后执行以下命令安装:

```
# ./configure
```

```
# make
```

```
# make install
```

安装完成后执行文件位于/usr/local/sbin, 如下所示:

```
# ls -l /usr/local/sbin
```

```
total 6480
```

```
-rwxr-xr-x 1 root root 1891727 Jan 31 23:37 bgpd
```

```
-rwxr-xr-x 1 root root 1090742 Jan 31 23:37 ospf6d
```

```
-rwxr-xr-x 1 root root 1398401 Jan 31 23:37 ospfd
```

```
-rwxr-xr-x 1 root root 824695 Jan 31 23:37 ripd
```

```
-rwxr-xr-x 1 root root 713611 Jan 31 23:37 ripngd
```

```
-rwxr-xr-x 1 root root 677048 Jan 31 23:37 zebra
```

配置样本文件位于/usr/local/etc, 如下所示:

```
# ls -l /usr/local/etc
```

```
total 44
```

```
-rw----- 1 root root 572 Jan 31 23:37 bgpd.conf.sample
```

```
-rw----- 1 root root 2801 Jan 31 23:37 bgpd.conf.sample2
```

```
-rw----- 1 root root 1110 Jan 31 23:37 ospf6d.conf.sample
```

```
-rw----- 1 root root 180 Jan 31 23:37 ospfd.conf.sample
```

```
-rw----- 1 root root 412 Jan 31 23:37 ripd.conf.sample
```

```
-rw----- 1 root root 396 Jan 31 23:37 ripngd.conf.sample
```

```
-rw----- 1 root root 375 Jan 31 23:37 zebra.conf.sample
```

三、运行

各个 zebra 模块运行时需要事先定义一个配置文件 (*.conf)，可直接使用样本配置文件。

```
# cd /usr/local/etc
```

```
# cp zebra.conf.sample zebra.conf
```

可通过-h 参数查看 help:

```
# zebra -h
```

```
Usage : zebra [OPTION...]
```

Daemon which manages kernel routing table management and redistribution between different routing protocols.

-b, --batch Runs in batch mode

-d, --daemon Runs in daemon mode

-f, --config_file Set configuration file name

-i, --pid_file Set process identifier file name

-k, --keep_kernel Don't delete old routes which installed by zebra.

-l, --log_mode Set verbose log mode flag

-A, --vty_addr Set vty's bind address

-P, --vty_port Set vty's port number

-r, --retain When program terminates, retain added route

by zebra.

-v, --version Print program version

-h, --help Display this help and exit

Report bugs to bug-zebra@gnu.org

可使用-d 参数以后台进程模式启动各模块，比如 zebra:

```
# zebra -d
```

查看 zebra 进程是否运行:

```
# ps -ef | grep zebra
```

```
root      3039      1  0 14:52 ?          00:00:00 zebra -d
```

```
root      3131  2963  0 15:31 pts/3    00:00:00 grep zebra
```

检查后台进程监听端口:

```
# lsof -i:2601
```

COMMAND	PID	USER	FD	TYPE	DEVICE	SIZE	NODE	NAME
---------	-----	------	----	------	--------	------	------	------

zebra	3039	root	10u	IPv6		9912		TCP
-------	------	------	-----	------	--	------	--	-----

*:discp-client (LISTEN)

zebra-0.95a 安装好后会自动往系统/etc/service 中添加定义:

```
# Ports numbered 2600 through 2606 are used by the zebra package without
# being registered. The primary names are the registered names, and the
# unregistered names used by zebra are listed as aliases.
hpstgmgr      2600/tcp      zebrasrv      # HPSTGMGR
hpstgmgr      2600/udp      zebrasrv      # HPSTGMGR
discp-client  2601/tcp      zebra         # discp client
discp-client  2601/udp      zebra         # discp client
discp-server  2602/tcp      ripd          # discp server
discp-server  2602/udp      ripd          # discp server
servicemeter  2603/tcp      ripngd        # Service Meter
servicemeter  2603/udp      ripngd        # Service Meter
nsc-ccs       2604/tcp      ospfd         # NSC CCS
nsc-ccs       2604/udp      ospfd         # NSC CCS
nsc-posa      2605/tcp      bgpd          # NSC POSA
nsc-posa      2605/udp      bgpd          # NSC POSA
netmon        2606/tcp      ospf6d        # Dell Netmon
netmon        2606/udp      ospf6d        # Dell Netmon
```

四、配置

zebra 是基本进程，诸如路由器 hostname、接口 IP 等基本信息都在这个模块中配置。

可通过以下命令进入 zebra 模块：

```
# telnet localhost 2601
```

```
Trying 127.0.0.1...
```

```
Connected to localhost.localdomain (127.0.0.1).
```

```
Escape character is '^]'.
```

```
Hello, this is zebra (version 0.95a).
```

```
Copyright 1996-2004 Kunihiro Ishiguro.
```

```
User Access Verification
```

```
Password:
```

密码缺省为 zebra，输入密码后进入到以下配置界面，是不是感觉进入到了一台真正的 Cisco 路由器？

```
[root@RHEL5 etc]# telnet localhost 2601
Trying 127.0.0.1...
Connected to localhost.localdomain (127.0.0.1).
Escape character is '^]'.

Hello, this is zebra (version 0.95a).
Copyright 1996-2004 Kunihiro Ishiguro.

User Access Verification

Password:
Router> en
Password:
Router#
```

就绪

操作模式及命令跟 Cisco IOS 很相似，如查看配置也是使用 show

running-config, 如下所示:

```
Router# sh run
```

Current configuration:

```
!
```

```
hostname Router
```

```
password zebra
```

```
enable password zebra
```

```
!
```

```
interface lo
```

```
!
```

```
interface eth1
```

```
ipv6 nd suppress-ra
```

```
!
```

```
interface peth0
```

```
    ipv6 nd suppress-ra
```

```
!
```

```
interface sit0
```

```
    ipv6 nd suppress-ra
```

```
!
```

```
interface vif0.0
```

```
    ipv6 nd suppress-ra
```

```
!
```

```
interface eth0

    ipv6 nd suppress-ra
!

interface vif0.1

    ipv6 nd suppress-ra
!

interface veth1

    ipv6 nd suppress-ra
!

interface vif0.2

    ipv6 nd suppress-ra
!

interface veth2

    ipv6 nd suppress-ra
!

interface vif0.3

    ipv6 nd suppress-ra
!

interface veth3

    ipv6 nd suppress-ra
!

interface xenbr0
```

```
ipv6 nd suppress-ra
```

```
!
```

```
line vty
```

```
!
```

```
end
```

下面是为接口 eth1 配置一个 ip 过程：

```
Router# conf t
Router(config)# int
Router(config)# interface eth1
Router(config-if)# ip address 192.168.11.8/24
    label      Label of this address
    secondary   Secondary IP address
    <cr>
Router(config-if)# ip address 192.168.11.8/24
Router(config-if)#
```

就绪

查看接口 eth1 配置：

```
Router# sh interface eth1
    <cr>
Router# sh interface eth1
Interface eth1
  index 2 metric 1 mtu 1500 <UP,BROADCAST,RUNNING,MULTICAST>
  HWaddr: 00:11:d8:ed:df:ef
  inet 192.168.11.8/24 broadcast 192.168.11.255
  inet6 fe80::211:d8ff:feed:dfef/64
    input packets 313, bytes 41266, dropped 0, multicast packets 0
    input errors 0, length 0, overrun 0, CRC 0, frame 0, fifo 0, missed 0
    output packets 255, bytes 22312, dropped 0
    output errors 0, aborted 0, carrier 0, fifo 0, heartbeat 0, window 0
    collisions 0
Router#
```

就绪

五、案例

案例网络拓扑如下：



说明：在 Redhat linux 服务器和 Cisco 2610 路由器上启用 RIP 动态路由器协议,实现个人 PC 机(10. 200. 51. 202/8)能访问到 Cisco 2610 路由器 Loopback 1 (172. 16. 1. 1/24)。

1、Cisco 2610 配置

```

!
hostname Router
!
enable password cisco
!
!
ip subnet-zero
!
!
interface Loopback1
 ip address 172.16.1.1 255.255.0.0
!
interface Ethernet0/0
 ip address 192.168.11.7 255.255.255.0
!
interface Serial0/0
 no ip address
 shutdown
!
router rip
 version 2
 network 172.16.0.0
 network 192.168.11.0
!
ip classless
no ip http server
!
!
line con 0
 transport input none
line aux 0
line vty 0 4
 password cisco
 login
!
end

```

2、Redhat linux 服务器配置：

确保 zebra 和 ripd 模块已运行：

```
# ps -ef | grep zebra
```

```
root          3039      1  0 14:52 ?           00:00:00 zebra -d
```

```
root          3314    2901  0 17:02 pts/1      00:00:00 grep zebra
```

```
# ps -ef | grep ripd
```

```
root          3052          1  0 14:54 ?          00:00:00 ripd -d
root          3316    2901   0 17:02 pts/1        00:00:00 grep ripd
```

配置 Redhat linux 服务器 Eth1 网卡 ip:

```
[root@RHEL5 ~]# telnet localhost 2601
Trying 127.0.0.1...
Connected to localhost.localdomain (127.0.0.1).
Escape character is '^]'.

Hello, this is zebra (version 0.95a).
Copyright 1996-2004 Kunihiro Ishiguro.

User Access Verification

Password:
Password:
Router> en
Password:
Router# conf t
Router(config)# int eth1
Router(config-if)# ip add 192.168.11.8/24
Router(config-if)# no shut
Router(config-if)# exit
Router(config)# exit
Router# copy runn
Router# copy running-config start
Router# copy running-config startup-config
Configuration saved to /usr/local/etc/zebra.conf
Router#
```

配置 Redhat linux 服务器 RIP 动态路由协议:

```
[root@RHEL5 ~]# telnet localhost 2602
Trying 127.0.0.1...
Connected to localhost.localdomain (127.0.0.1).
Escape character is '^]'.

Hello, this is zebra (version 0.95a).
Copyright 1996-2004 Kunihiro Ishiguro.

User Access Verification

Password:
ripd> enable
ripd# conf t
ripd(config)# router rip
ripd(config-router)# version 2
ripd(config-router)# network 10.0.0.0/8
ripd(config-router)# network 192.168.11.0/24
ripd(config-router)# exit
ripd(config)# exit
ripd# copy run star
Configuration saved to /usr/local/etc/ripd.conf
ripd#
```

3、查看

在 Cisco 2610 路由器检查 RIP 是否生效：

```
Router#sh ip route
Codes: C - connected, S - static, I - IGRP, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2, E - EGP
       i - IS-IS, L1 - IS-IS level-1, L2 - IS-IS level-2, ia - IS-IS inter area
       * - candidate default, U - per-user static route, o - ODR
       P - periodic downloaded static route

Gateway of last resort is not set

C       172.16.0.0/16 is directly connected, Loopback1
C       192.168.11.0/24 is directly connected, Ethernet0/0
R       10.0.0.0/8 [120/1] via 192.168.11.8, 00:00:03, Ethernet0/0
Router#
```

从 sh ip route 输出结果（白色加显部分）可知，Cisco 2610 路由器已成功学习到 Redhat linux 服务器广播过来的 RIP 路由条目。

在 Redhat linux 服务器检查 RIP 是否生效：

```
[root@RHEL5 ~]# telnet localhost 2602
Trying 127.0.0.1...
Connected to localhost.localdomain (127.0.0.1).
Escape character is '^]'.

Hello, this is zebra (version 0.95a).
Copyright 1996-2004 Kunihiro Ishiguro.

User Access Verification

Password:
ripd> en
ripd# sh ip rip
Codes: R - RIP, C - connected, O - OSPF, B - BGP
       (n) - normal, (s) - static, (d) - default, (r) - redistribute,
       (i) - interface

   Network        Next Hop        Metric From        Time
C(i) 10.0.0.0/8    0.0.0.0          1 self
R(n) 172.16.0.0/16 192.168.11.7      2 192.168.11.7     02:50
C(i) 192.168.11.0/24 0.0.0.0          1 self
ripd#
```

从 sh ip rip 输出结果（白色加显部分）可知，Redhat linux 服务器已成功学习到 Cisco 2610 路由器广播过来的 RIP 路由条目。

4、测试

在个人 PC 机上添添加一条到 172.16.1.1/32 的静态路由：

```
C:\>netstat -r

Route Table
=====
Interface List
0x1 ..... MS TCP Loopback interface
0x2 ...00 17 31 fb 7a b7 ..... Realtek RTL8139/810x Family Fast Ethernet NIC -
Deterministic Network Enhancer Miniport
0x3 ...00 ac 39 38 f7 7b ..... SoftEther Virtual LAN Card Adapter - Determinist
ic Network Enhancer Miniport
=====
Active Routes:
Network Destination        Netmask          Gateway             Interface           Metric
0.0.0.0                    0.0.0.0          10.200.49.254        10.200.51.202         1
10.0.0.0                    255.0.0.0         10.200.51.202        10.200.51.202        20
10.200.51.202              255.255.255.255    127.0.0.1           127.0.0.1            20
10.243.67.39               255.255.255.255    10.200.6.200         10.200.51.202         1
10.255.255.255             255.255.255.255    10.200.51.202        10.200.51.202        20
127.0.0.0                  255.0.0.0         127.0.0.1           127.0.0.1            1
172.16.1.1                 255.255.255.255    10.200.2.216         10.200.51.202         1
172.20.0.0                 255.255.0.0        10.200.6.201         10.200.51.202         1
```

从个人 PC 机通过 ping/tracert/telnet 测试网络：

```
C:\>ping 172.16.1.1

Pinging 172.16.1.1 with 32 bytes of data:

Reply from 172.16.1.1: bytes=32 time=1ms TTL=254
Reply from 172.16.1.1: bytes=32 time=1ms TTL=254
Reply from 172.16.1.1: bytes=32 time=1ms TTL=254
Reply from 172.16.1.1: bytes=32 time=1ms TTL=254

Ping statistics for 172.16.1.1:
    Packets: Sent = 4, Received = 4, Lost = 0 (0% loss),
    Approximate round trip times in milli-seconds:
        Minimum = 1ms, Maximum = 1ms, Average = 1ms

C:\>tracert 172.16.1.1

Tracing route to 172.16.1.1 over a maximum of 30 hops

  1    <1 ms    <1 ms    <1 ms    rhel5 [10.200.2.216]
  2     1 ms     1 ms     1 ms     172.16.1.1

Trace complete.

C:\>_
```

```
C:\> Telnet 172.16.1.1
```

```
User Access Verification
Password:
```

从以上测试可知，Redhat linux 服务器运行的 RIP 动态路由协议能成功和 Cisco 2610 路由器运行的 RIP 动态路由协议配合工作。

5、特别说明

Redhat linux 服务器必须启用路由转发功能：

```
# echo "1" > /proc/sys/net/ipv4/ip_forward
```

六、结尾语

- 1、对于小型公司来说，可能没这么多资金购买硬件路由产品，低成本的 PC 机+linux+zebra 提供了一个可行的解决方案；
- 2、现在很多培训机构，特别是小型培训机构，为了节省投资，也会采取这种架构搭建实验环境。

E-Mail: wzknet@hotmail.com

Blog: <http://k968888.blog.sohu.com>

2008 年 2 月 2 日 广州